

# CS536 Lab 3

Borna Tavasoli

October 15, 2025

we don't make mistakes, we have happy  
accidents

---

*Bob Ross*

## Problem 1.

### 1.1 System Set-Up, Traffic Generation, and Capture

To begin, the Ethernet interface `veth0` on one of the amber machines in the HAAS G050 lab was used. Promiscuous mode was required for sniffing Ethernet frames, which necessitated superuser privileges.

```
sudo /usr/local/etc/tcpdumpwrap -veth0 -c 32 -w > etherlogfile
```

This command captures 32 Ethernet frames and stores them into `etherlogfile`. The captured file was placed under the directory `lab3/v1/` as required.

**Traffic generation:** Traffic was generated using the UDP ping application from Lab 2. The ping server was bound to IP address `192.168.1.1`, and the client was executed from the same machine with:

```
veth 'udppingc 192.168.1.1 8080 123456 8081 100'
```

Here, the `veth` wrapper executed the client on the virtual interface with IP `192.168.1.2`. This allowed packet exchange between `192.168.1.1` and `192.168.1.2` over a virtual Ethernet link. The configuration of `veth0` at the remote end was verified using:

```
veth 'ifconfig veth0'
```

### 1.2 Traffic Analysis

After generating sufficient traffic, the file `etherlogfile` was analyzed using Wireshark.

**Step 1 - Opening in Wireshark:** The file was opened in Wireshark with:

```
/usr/bin/wireshark etherlogfile
```

Wireshark automatically decoded the Ethernet, IP, and UDP headers, displaying packet structure. Results is shown in Figure 1 and 2.

[h]

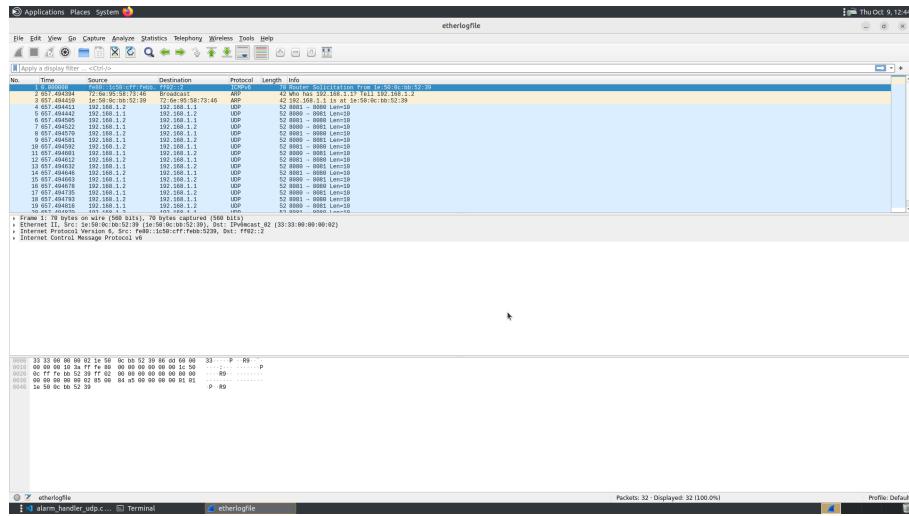


Figure 1: Wireshark environment

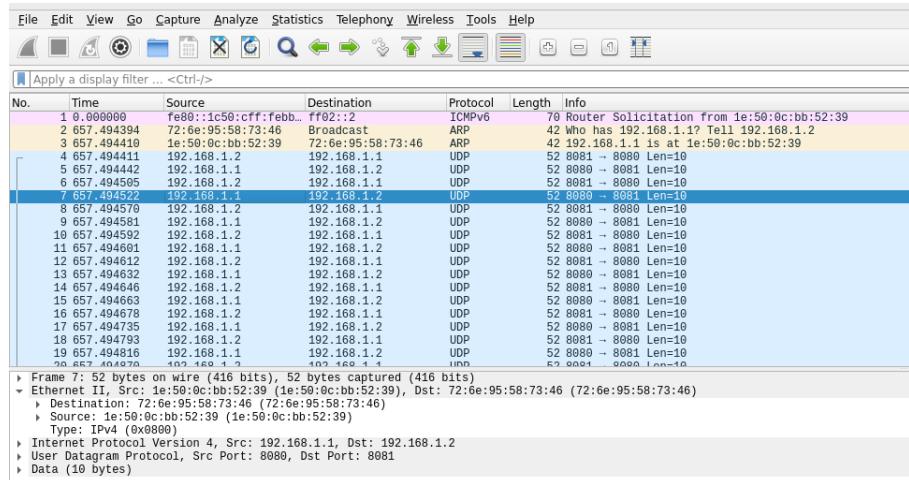


Figure 2: Wireshark closer look

**Step 2 - Identifying MAC addresses:** The MAC addresses corresponding to 192.168.1.1 and 192.168.1.2 were obtained using:

```
ifconfig -a
veth 'ifconfig -a'
```

These were then matched with the source and destination addresses in the Ethernet frames. Results is shown in Figure 3 and 4.

Note that 0xc0a80101 is hex for 192.168.1.1. Similarly 0xc0a80102 is hex for 192.168.1.2.

**Step 3 - Checking Frame Type:** The Ethernet Type field was inspected and confirmed to correspond to IPv4 (0x0800), meaning the frames used the DIX (Ethernet II) format. Results is shown in Figure 5.

**Step 4 - Inspecting IP and UDP headers:** The source and destination IP addresses, as well as port numbers, were confirmed to match the ones used in the udppingc application. Results is shown in Figure 6.

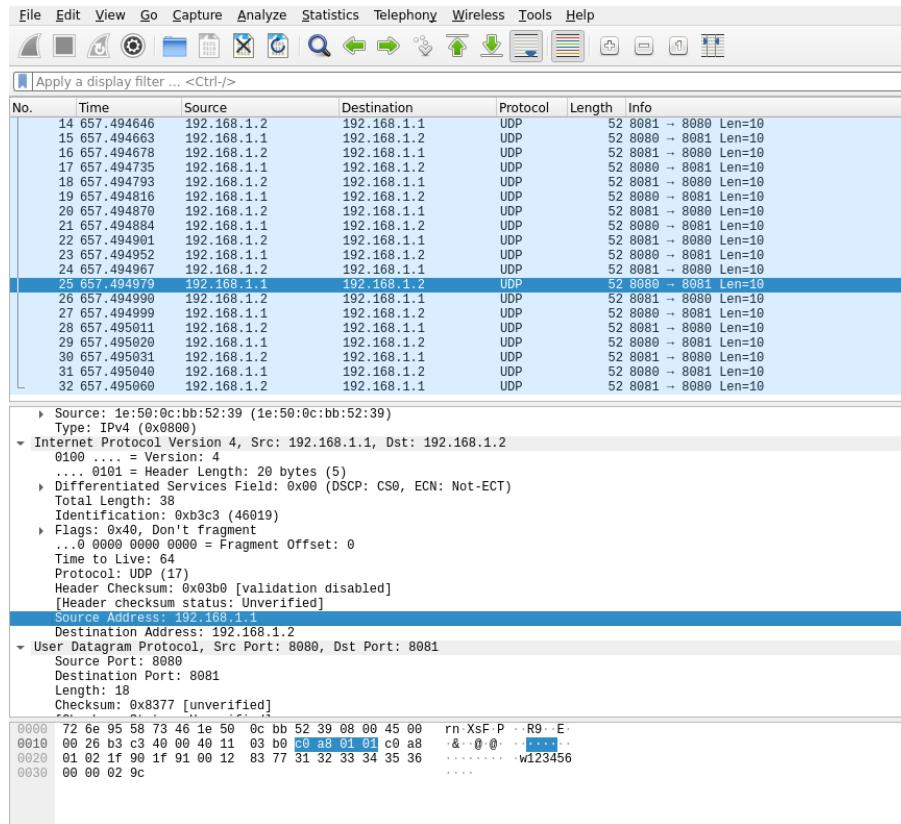


Figure 3: Source Mac Address

**Step 5 - Inspecting application layer payload:** The remaining bytes beyond the UDP header correspond to the application-layer payload sent via `sendto()`. By switching to “View → Packet Bytes” in Wireshark, the payload was viewed in raw hexadecimal format to verify transmitted data. Results is shown in Figure 7.

We can verify that this payload is consistent among packets. Results is shown in Figure 8.

**Verification with tcpdump:** To double-check, the following command was used:

```
tcpdump -r - < etherlogfile
```

This confirmed consistent IP addresses, port numbers, and packet structure as seen in Wireshark. Results is shown in Figure 9.

## Summary

- The Ethernet frames were confirmed to use Ethernet II (DIX) format with Type field value 0x0800.
- IP headers contained source and destination IPs of 192.168.1.1 and 192.168.1.2.
- UDP headers contained the correct source (8080) and destination (8081) port numbers used by the client and server.
- The application payload matched the data generated by the `udppingc` ping client.
- Both Wireshark and tcpdump produced consistent interpretations of the captured data.

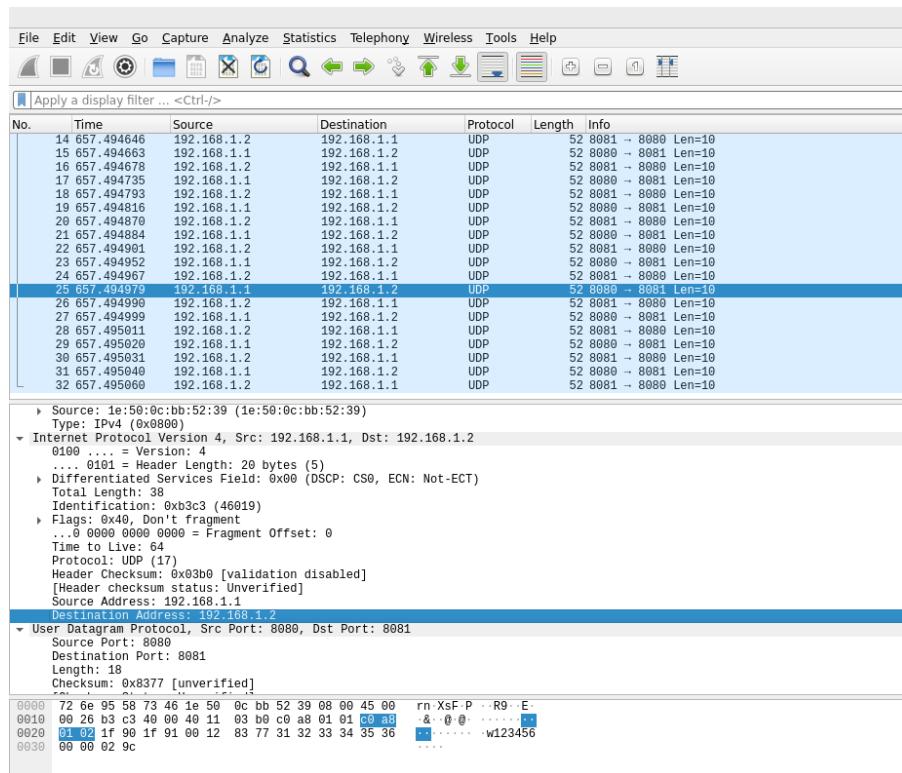


Figure 4: Destination Mac Address

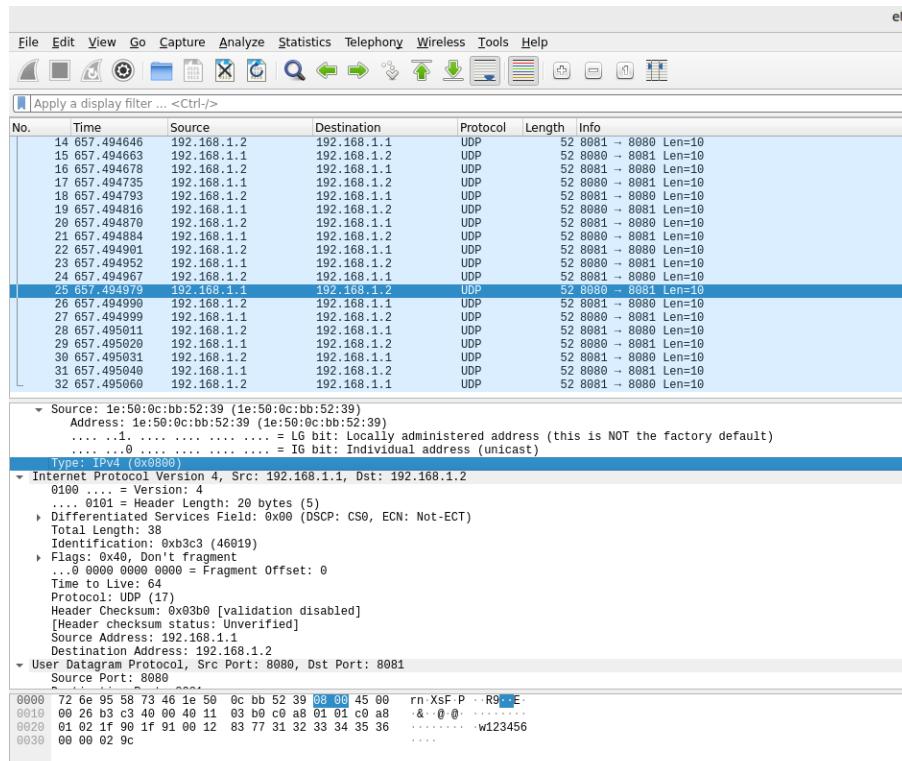


Figure 5: Ethernet Type

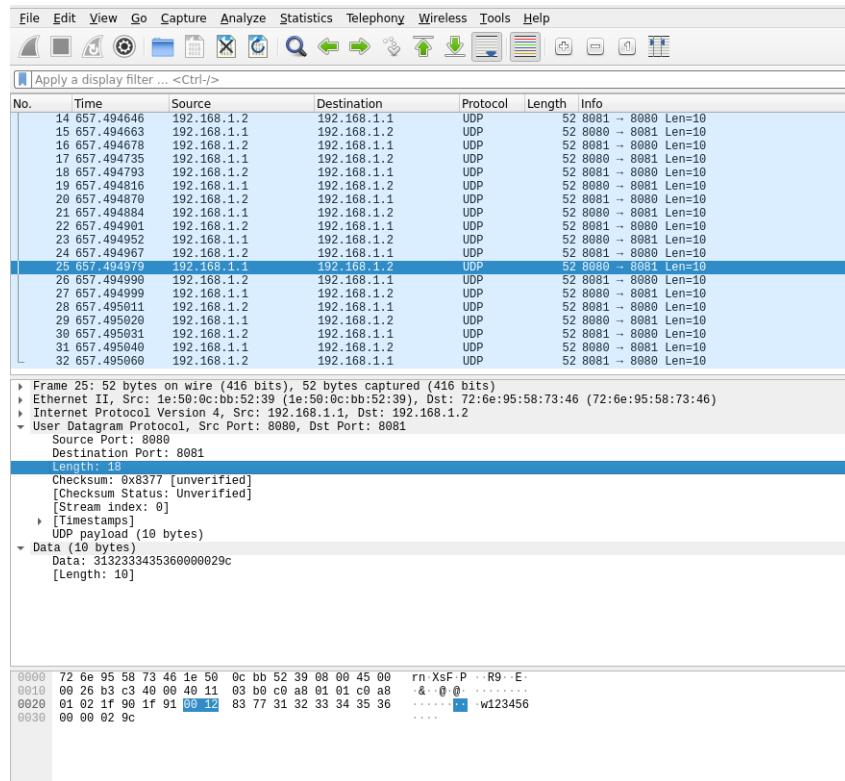


Figure 6: UDP Headers

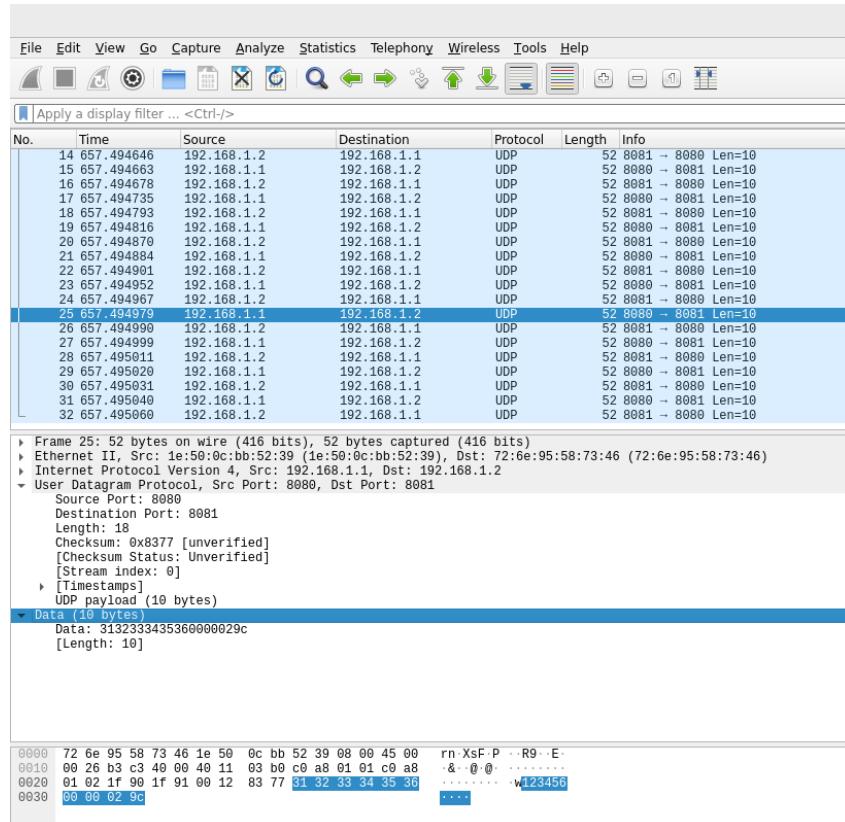


Figure 7: Payload View

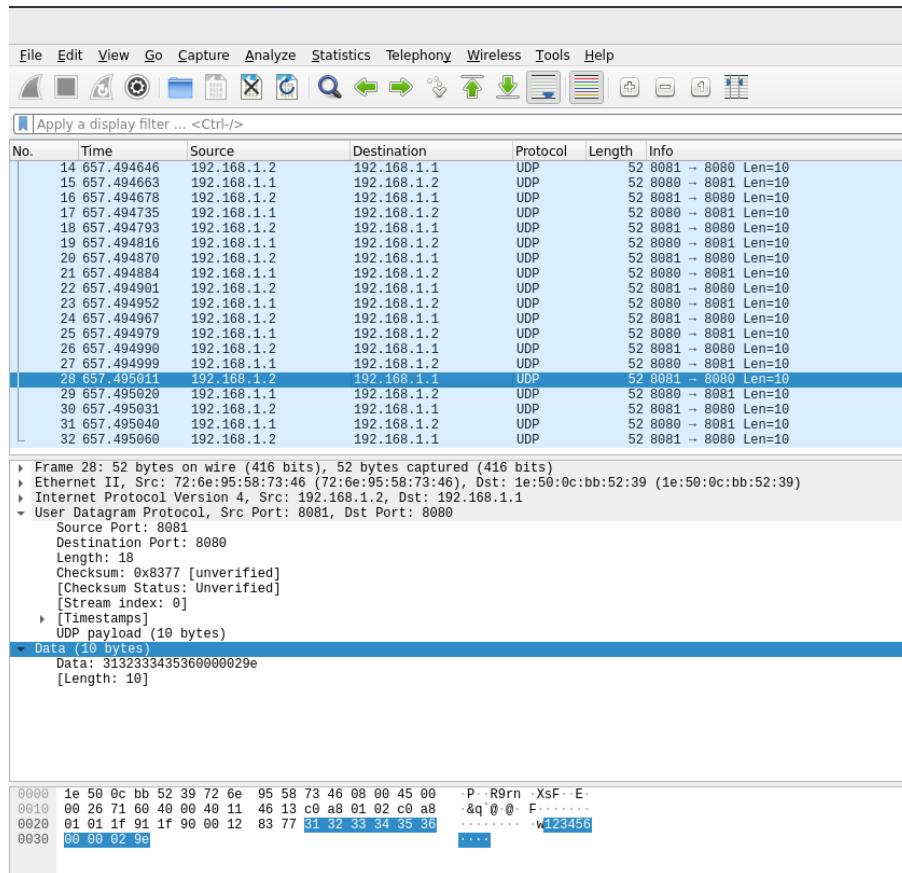


Figure 8: Checking payload consistency

```

data 57 $ tcpdump -r - < etherlogfile
reading from file -, link-type EN10MB (Ethernet), snapshot length 262144
12:00:13.240318 IP6 fe80::1c50:cff:febb:5239 > ip6-allrouters: ICMP6, router solicitation, length 16
12:11:10.734712 ARP, Request who-has 192.168.1.1 tell 192.168.1.2, length 28
12:11:10.734728 ARP, Reply 192.168.1.1 is-at 1e:50:0c:bb:52:39 (oui Unknown), length 28
12:11:10.734729 IP 192.168.1.2.8081 > 192.168.1.1.8080: UDP, length 10
12:11:10.734760 IP 192.168.1.1.8080 > 192.168.1.2.8081: UDP, length 10
12:11:10.734823 IP 192.168.1.2.8081 > 192.168.1.1.8080: UDP, length 10
12:11:10.734840 IP 192.168.1.1.8080 > 192.168.1.2.8081: UDP, length 10
12:11:10.734888 IP 192.168.1.2.8081 > 192.168.1.1.8080: UDP, length 10
12:11:10.734899 IP 192.168.1.1.8080 > 192.168.1.2.8081: UDP, length 10
12:11:10.734910 IP 192.168.1.2.8081 > 192.168.1.1.8080: UDP, length 10
12:11:10.734919 IP 192.168.1.1.8080 > 192.168.1.2.8081: UDP, length 10
12:11:10.734930 IP 192.168.1.2.8081 > 192.168.1.1.8080: UDP, length 10
12:11:10.734950 IP 192.168.1.1.8080 > 192.168.1.2.8081: UDP, length 10
12:11:10.734964 IP 192.168.1.2.8081 > 192.168.1.1.8080: UDP, length 10
12:11:10.734981 IP 192.168.1.1.8080 > 192.168.1.2.8081: UDP, length 10
12:11:10.734996 IP 192.168.1.2.8081 > 192.168.1.1.8080: UDP, length 10
12:11:10.735053 IP 192.168.1.1.8080 > 192.168.1.2.8081: UDP, length 10
12:11:10.735111 IP 192.168.1.2.8081 > 192.168.1.1.8080: UDP, length 10
12:11:10.735134 IP 192.168.1.1.8080 > 192.168.1.2.8081: UDP, length 10
12:11:10.735188 IP 192.168.1.2.8081 > 192.168.1.1.8080: UDP, length 10
12:11:10.735202 IP 192.168.1.1.8080 > 192.168.1.2.8081: UDP, length 10
12:11:10.735219 IP 192.168.1.2.8081 > 192.168.1.1.8080: UDP, length 10
12:11:10.735270 IP 192.168.1.1.8080 > 192.168.1.2.8081: UDP, length 10
12:11:10.735285 IP 192.168.1.2.8081 > 192.168.1.1.8080: UDP, length 10
12:11:10.735297 IP 192.168.1.1.8080 > 192.168.1.2.8081: UDP, length 10
12:11:10.735308 IP 192.168.1.2.8081 > 192.168.1.1.8080: UDP, length 10
12:11:10.735317 IP 192.168.1.1.8080 > 192.168.1.2.8081: UDP, length 10
12:11:10.735329 IP 192.168.1.2.8081 > 192.168.1.1.8080: UDP, length 10
12:11:10.735338 IP 192.168.1.1.8080 > 192.168.1.2.8081: UDP, length 10
12:11:10.735349 IP 192.168.1.2.8081 > 192.168.1.1.8080: UDP, length 10
12:11:10.735358 IP 192.168.1.1.8080 > 192.168.1.2.8081: UDP, length 10
12:11:10.735378 IP 192.168.1.2.8081 > 192.168.1.1.8080: UDP, length 10
data 58 $ []

```

Figure 9: TCP dump output

## Problem 2.

### Asynchronous Concurrency Management

The sender implementation introduces asynchrony through the `SIGIO` signal handler, which processes incoming ACK packets while the main synchronous loop transmits data packets. This creates potential race conditions when both the signal handler and main loop access shared data structures, specifically the `ack_received[]` array that tracks which packets have been successfully acknowledged.

Our implementation uses a reader-writer separation strategy that avoids the need for complex synchronization primitives such as mutexes or signal masking. If you look at the code:

- Synchronous code (main loop): Only reads from the `ack_received[]` array
- Asynchronous code (SIGIO handler): Only writes to the `ack_received[]` array

By doing so, we avoid race conditions because:

1. Neither component performs atomic read-modify-write sequences on shared data
2. The signal handler only transitions flags from 0 to 1, never reverses them
3. If the main loop reads a stale value (0 instead of 1), it simply retransmits the packet unnecessarily, which is safe because the receiver handles duplicates correctly
4. On modern architectures, byte-sized writes are atomic which means we won't have any incomplete reads of individual array elements

Thus by enforcing strict read-only and write-only roles for the synchronous and asynchronous components respectively, our implementation achieves correct concurrent operation without the overhead and complexity of explicit synchronization primitives.

### Duplicate File Transfer Initiation Packet Handling

After the receiver sends an acknowledgment to a file transfer initiation packet, it sets a 250ms timer to wait for the first data packet. During this waiting period, duplicate initiation packets may arrive from the sender: either due to network delays causing the sender to timeout and retransmit, or due to genuine packet duplication in the network. The receiver must handle these duplicates correctly without disrupting the established session.

In our approach, the receiver maintains a `transfer_active` flag that tracks whether a file transfer session is currently in progress. This flag transitions through three states:

1. Idle (`transfer_active = 0`): Waiting for a new initiation packet
2. Handshake (`transfer_active = 1, next_expected = 0`): Init ACK sent, waiting for first data packet
3. Active transfer (`transfer_active = 1, next_expected > 0`): Receiving data packets

If a duplicate initiation packet (16 bytes) arrives before the first data packet, the receiver responds with an idempotent acknowledgment and resends the 6-byte filename ACK without modifying any state (memory buffers, tracking arrays, session parameters remain unchanged). This is safe because multiple identical ACKs don't corrupt protocol state. Moreover, if the original ACK was lost, the duplicate provides recovery. We will then restart the 250ms timer automatically and wait for the data. Once the first data packet arrives, initiation packets from different senders are ignored to prevent session hijacking. Note that once the connection is established, getting 16 byte packets (whole packet or remainder of a packet) will not cause any issues since we will be in the third stage.

## Performance Optimization

### Experimental Setup and Result

We conducted performance experiments using a 102.0 kB test file to evaluate the impact of protocol parameters on transfer completion time. The two primary tunable parameters are:

- Micropace: Delay in microseconds between consecutive packet transmissions
- Payload size: Number of data bytes per UDP packet (excluding 4-byte sequence number header)

All tests were performed in the local lab environment (low latency, high bandwidth, minimal background loss). Below is our result:

Micropace (μs)	Payload (bytes)	Transfer Time (ms)
500	200	100
500	500	35
500	1000	22
500	2000	14
500	3000	10
300	1000	26
100	1000	15

Table 1: Transfer completion time for 102.0 kB file

### Analysis

**Impact of Payload Size** Holding micropace constant at 500μs while varying payload size reveals a clear inverse relationship between payload size and transfer time. Key observations:

- Larger payloads reduce the number of packets required: 102 kB with 200-byte payload requires 510 packets, while 3000-byte payload requires only 34 packets
- Fewer packets mean less per-packet overhead (UDP/IP headers, system calls, protocol bookkeeping)
- Transfer time improves dramatically: from 100ms (200B payload) to 10ms (3000B payload)

However, payload size is constrained by the maximum UDP payload that avoids IP fragmentation. We limited our experiments to payloads  $\leq$  3000 bytes to stay safely below the typical MTU of 1500 bytes.

**Impact of Micropace** Holding payload constant at 1000 bytes while varying micropace shows diminishing returns:

- Reducing micropace from 500μs to 100μs improves time from 22ms to 15ms (32% improvement)
- However, reducing micropace from 500μs to 300μs *increases* time from 22ms to 26ms

The non-monotonic behavior at 300μs suggests measurement variability or system scheduling effects at small time scales. The general trend shows that lower micropace (faster transmission rate) reduces completion time by minimizing idle time between packet transmissions.

Through iterative experimentation, we identified that a configuration of approximately:

- **Micropace:** 300 $\mu$ s
- **Payload:** 8000 bytes

yields transfer times around **3ms** for the 102 kB file in our lab environment. This is a  $33 \times$  speedup compared to the baseline configuration (500 $\mu$ s, 200B payload).

### Final Notes

While decreasing micropace improves performance in ideal conditions, excessively low values introduce reliability concerns:

- Bursts of packets may exceed the receiver's socket buffer capacity, causing kernel-level packet drops
- In shared network environments, aggressive sending can overwhelm switch buffers or compete unfairly with other traffic
- Higher packet loss rates trigger more retransmissions, potentially negating performance gains

We observed that micropace values below 100 $\mu$ s occasionally resulted in failed transfers due to excessive packet loss, even in the controlled lab environment.

While larger payloads improve efficiency, they have drawbacks:

- Larger payloads mean more wasted bandwidth when a packet is lost and must be retransmitted
- For small files, large payloads may not improve performance and can increase latency variance

Performance optimization in reliable UDP protocols involves balancing throughput (achieved through large payloads and low micropace) against reliability (which degrades under aggressive parameters). For small file transfers in low-loss, low-latency environments like our lab, the protocol can achieve sub-5ms completion times with careful parameter tuning. However, these optimal parameters are environment-specific and may require adjustment for production deployments with higher loss rates, variable latency, or shared network resources.

### Bonus

**File:** `m2_ch1_2024-02-19.pcap`

**Selected interval:** 13:00:05.07 – 13:00:05.60 (first 0.5 s, 806 frames)

- **Network type:** 802.11b/g/n (2.4 GHz band, Channel 1).
- **Frame types observed:** Predominantly *Beacon* frames, with occasional *Probe Request*, *Acknowledgment*, *Clear-to-Send (CTS)*, and *Null Function* frames.
- **Data rates:** Initially 1 Mbps for the first few frames, later increasing up to 12 Mbps.
- **SSID(s):** PAL3.0, PAL-Recreational, attwifi, eduroam
- **Signal strength  $\approx -65$  dBm**
- **Signal strength ratio  $\approx [-5, 90]$  dBm**
- **Notes:** The presence of multiple beacon frames suggests nearby access points broadcasting their presence. The mixture of management and control frames indicates active association and channel scanning activity.

Below is a summary of our result:

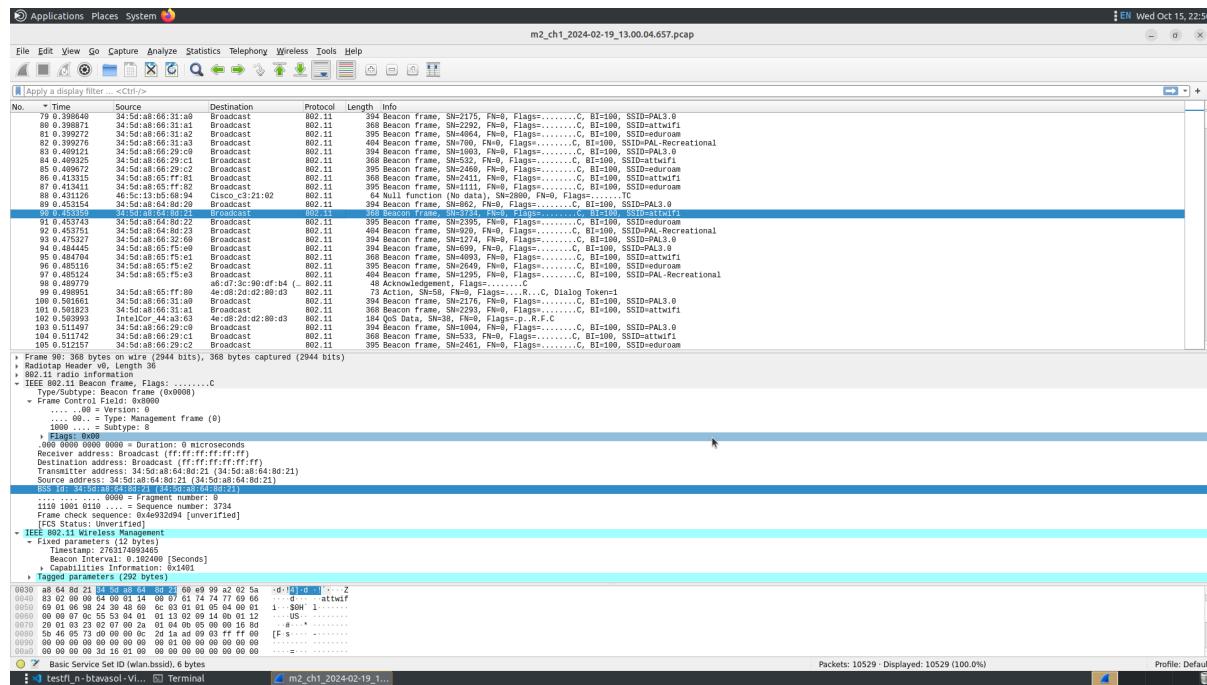


Figure 10: Bssid and ssid

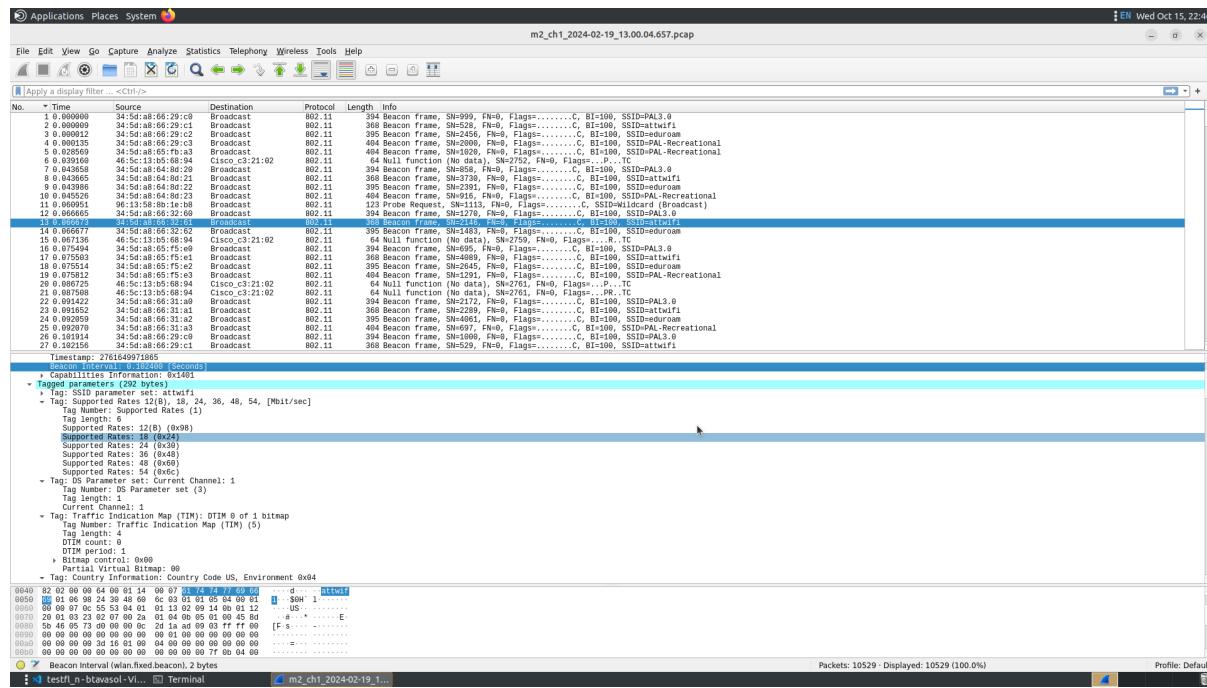


Figure 11: Channel info

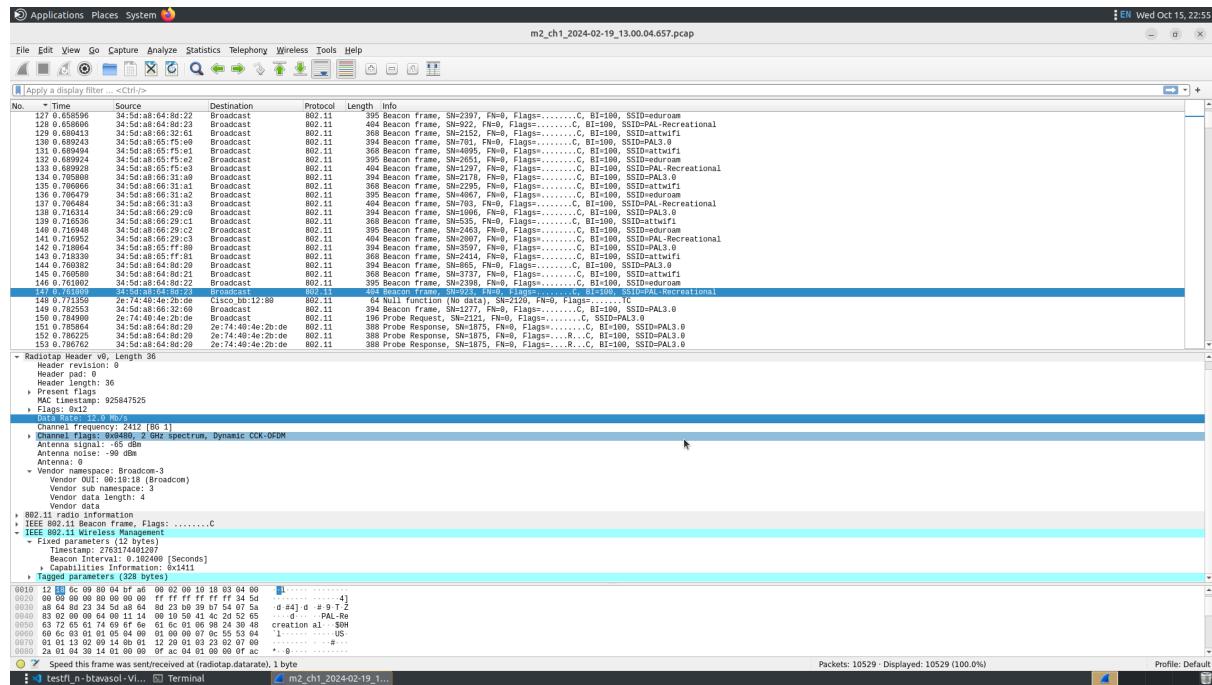


Figure 12: Data rate info

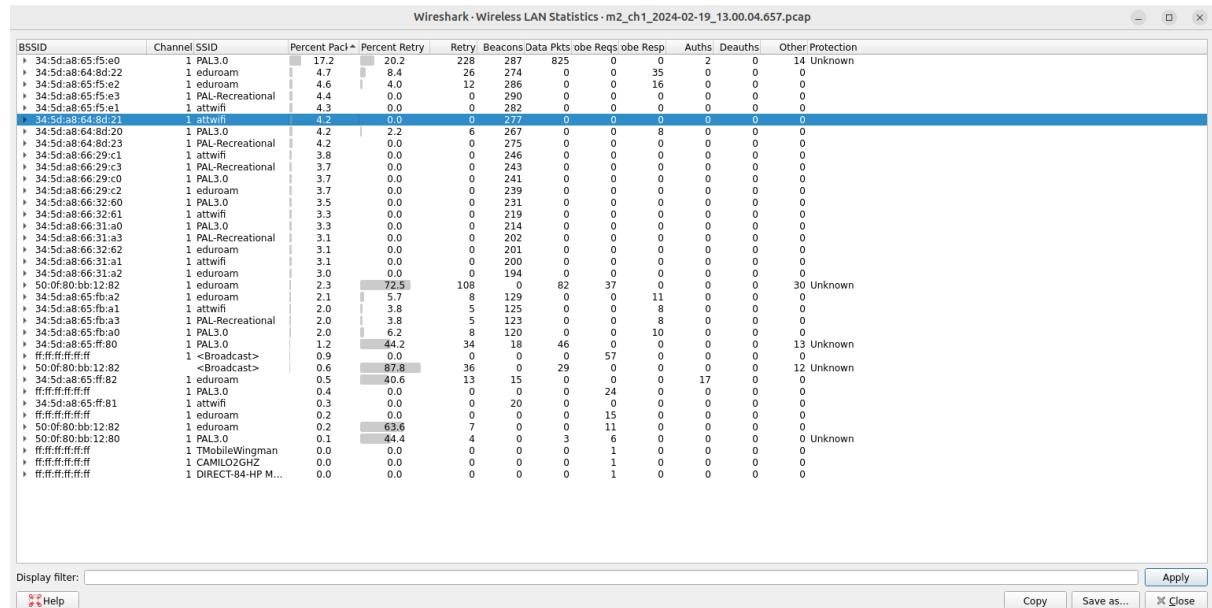


Figure 13: Frame type summary

```

187 0 7.065484 34:5d:8f:66:31:a3 Broadcast 802.11 484 Beacon frame, SH=763, FN=0, Flags=.....C, BI=100, SSID=PAL_Recreational
188 0 7.16314 34:5d:8f:66:29:c9 Broadcast 802.11 394 Beacon frame, SH=1006, FN=0, Flags=.....C, BI=100, SSID=PAL_0
139 0 7.16536 34:5d:8f:66:29:c1 Broadcast 802.11 368 Beacon frame, SH=53, FN=0, Flags=.....C, BI=100, SSID=attwifi
140 0 7.16948 34:5d:8f:66:29:c2 Broadcast 802.11 395 Beacon frame, SH=2463, FN=0, Flags=.....C, BI=100, SSID=eduroam
141 0 7.16957 34:5d:8f:66:29:c3 Broadcast 802.11 404 Beacon frame, SH=2007, FN=0, Flags=.....C, BI=100, SSID=PAL_Recreational
142 0 7.18064 34:5d:8f:65:ff:80 Broadcast 802.11 394 Beacon frame, SH=3597, FN=0, Flags=.....C, BI=100, SSID=PAL3.0
143 0 7.18330 34:5d:8f:64:ff:81 Broadcast 802.11 368 Beacon frame, SH=2414, FN=0, Flags=.....C, BI=100, SSID=attwifi
144 0 7.669382 34:5d:8f:64:8d:11 Broadcast 802.11 394 Beacon frame, SH=665, FN=0, Flags=.....C, BI=100, SSID=PAL3.0
145 0 7.66558 34:5d:8f:64:8d:12 Broadcast 802.11 368 Beacon frame, SH=3737, FN=0, Flags=.....C, BI=100, SSID=attwifi
146 0 7.66562 34:5d:8f:64:8d:13 Broadcast 802.11 395 Beacon frame, SH=1000, FN=0, Flags=.....C, BI=100, SSID=eduroam
147 0 7.66563 34:5d:8f:64:8d:14 Broadcast 802.11 404 Beacon frame, SH=3599, FN=0, Flags=.....C, BI=100, SSID=PAL_Recreational

148 0 7.733598 77:74:48:4e:b2:de Cisco_b2:12:80 802.11 64 Null function (No data), SH=2120, FN=0, Flags=.....TC
149 0 7.82553 34:5d:8f:66:32:60 Broadcast 802.11 394 Beacon frame, SH=1277, FN=0, Flags=.....C, BI=100, SSID=PAL3.0
150 0 7.849090 77:74:48:4e:b2:de Broadcast 802.11 196 Probe Request, SH=2121, FN=0, Flags=.....C, SSID=PAL3.0
151 0 7.85864 34:5d:8f:64:8d:20 2e:74:40:4e:2b:de 802.11 388 Probe Response, SH=1875, FN=0, Flags=.....C, BI=100, SSID=PAL3.0
152 0 7.86762 34:5d:8f:64:8d:21 2e:74:40:4e:2b:de 802.11 388 Probe Response, SH=1875, FN=0, Flags=.....R..C, BI=100, SSID=PAL3.0
153 0 7.86762 34:5d:8f:64:8d:20 2e:74:40:4e:2b:de 802.11 388 Probe Response, SH=1875, FN=0, Flags=.....R..C, BI=100, SSID=PAL3.0

MAC timestamp: 925847525
+ Flags: 0x10
Data rate: 12.0 Mb/s
Channel frequency: 2412 [BG 1]
+ Channel flags: 0x0480, 2 GHz spectrum, Dynamic CCK-OFDM
Antenna signal: -65 dBm
Antenna noise: -90 dBm
Antenna: 0
+ Vendor namespace: Broadcom-3
Vendor QID: 00:10:18:18 (Broadcom)
Vendor sub namespace: 3
Vendor data length: 4
Vendor data
+ 802.11 radio information
PHY type: 802.11g (ERP) (6)
Proprietary mode: None (0)
Data rate: 12.0 Mb/s
Channel: 1
Frequency: 2412MHz
Signal/noise ratio (dB): -65 dBm
Noise level (dBm): -90 dBm
Signal/noise ratio (dB): 25 dB
TSF timestamp: 925847525
+ [Duration: 268us]
+ IEEE 802.11 Beacon frame, Flags: .....
+ IEEE 802.11 Wireless Management

0020 00 00 00 00 00 00 ff ff ff ff ff ff ff ff 34 5d .....4]
0030 a8 64 80 23 34 50 a8 64 8d 23 b0 39 b7 5d 07 5a [d #4] d # 9 T Z
0040 83 02 00 00 64 00 11 14 00 10 50 41 4c 24 52 65 .....d ... - PAL_Re
0050 63 66 00 00 63 00 01 01 00 01 01 01 01 01 01 01 creation al _S0H
0060 63 66 03 01 04 05 04 00 01 02 03 04 05 55 53 04 1 .. .0S
0070 01 01 13 02 09 14 0b 01 12 20 01 03 23 02 07 09 .....#
0080 2a 01 04 30 14 01 00 09 08 ac 04 01 00 00 00 09 ac * 0 .....-
0090 04 01 00 00 09 ac 01 28 09 36 03 c9 b7 00 05 .....( - .....
```

Figure 14: Signal strength

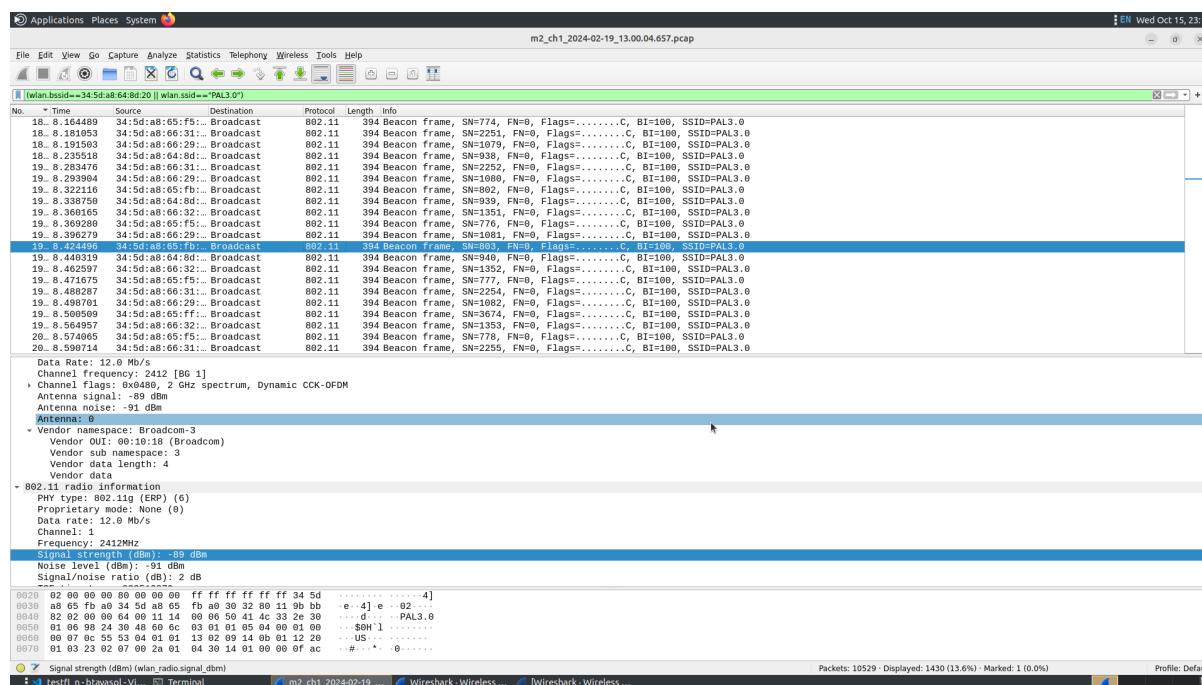


Figure 15: SNR of hightraffic basic service sets and their MAC addresses

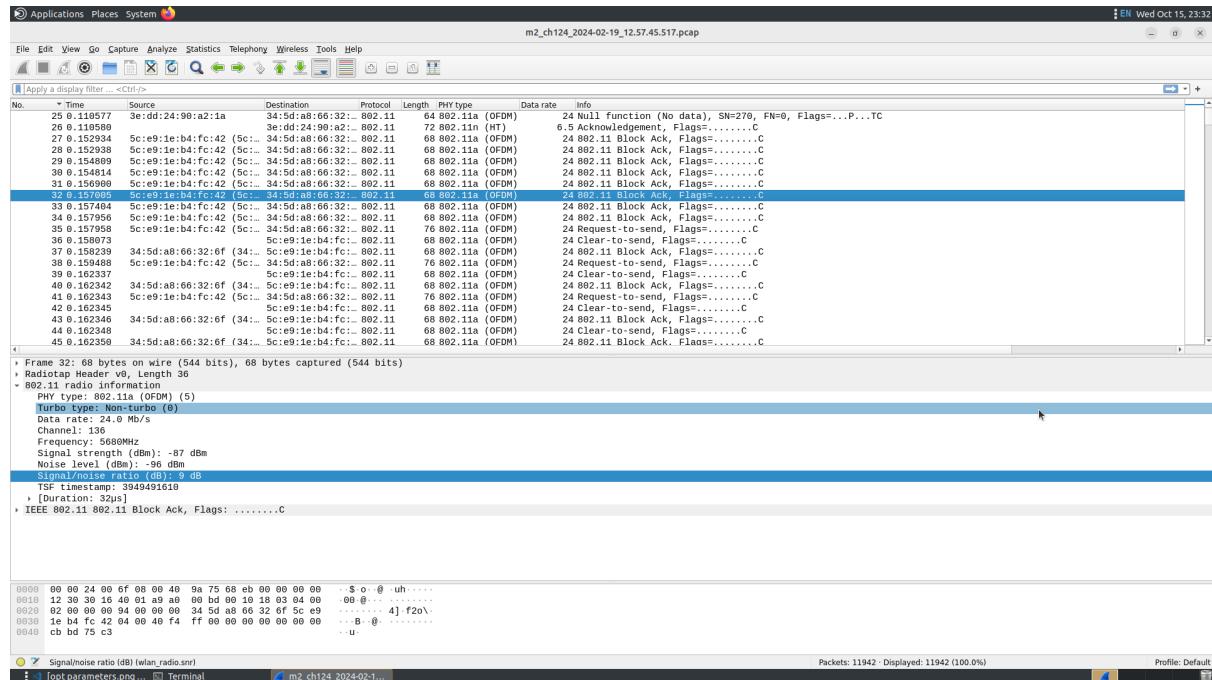


Figure 16: SNR, network type, and signal strength

**File: m2\_ch124\_2024-02-19.pcap****Selected interval:** [Specify e.g., 12:57:45.88 – 12:57:46.39] (about 256 packets)

- Network type:** 802.11a (5 GHz band, Channel 136).
- Frame types observed:** Mostly Block Ack, RTS and CTS, some beacon
- Data rates:** 24 to 12 Mb/s (we also have 6, 6.5, and 48)
- SSID(s):** Same as before
- Signal strength:**  $\approx -80$  dBm
- Signal strength ratio:**  $\approx [9, 15]$  dBm
- Notes:** Signal seems to be stronger than before. Also we have a wider band.

Below is a summary of our results:

**File: m2\_ch153\_2024-02-19.pcap****Selected interval:** The complete interval.

- Network type:** Likely 802.11a same as before (5 GHz band, Channel 153).
- Frame types observed:** Block Ack, RTS, CTS, VHT/HE NDP Announcement, Probe Request
- Data rates:** 6, 6.5 and 24
- SSID(s):** Wildcard, PAL3.0, Hilton Honors, Rajkiran, FreeORDWiFi, SSN, Hotel bobo

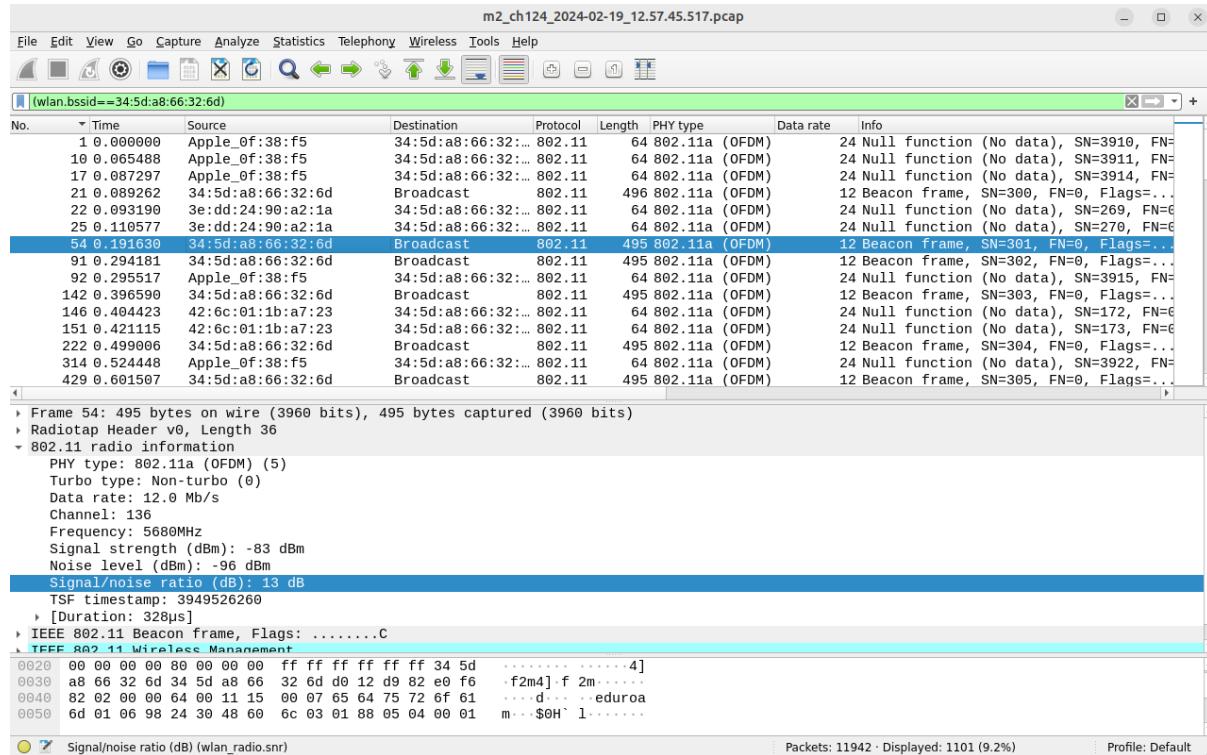


Figure 17: SNR of hightraffic basic service sets and their MAC addresses

- **Signal strength: -86 dBm**
- **Signal strength ratio: 45 dBm**
- **Notes:** Same band as before. Better signal ratio. Less different data rates. Different SSIDs.

Below is a summary of our findings:

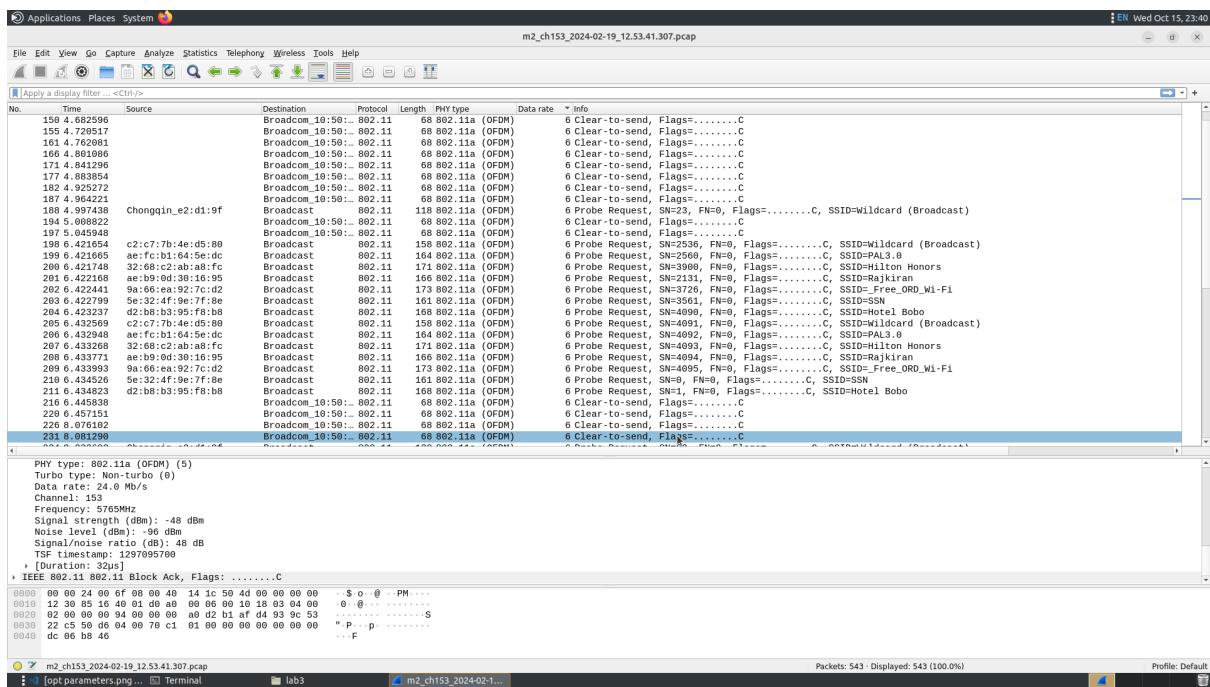


Figure 18: SSID and datarates

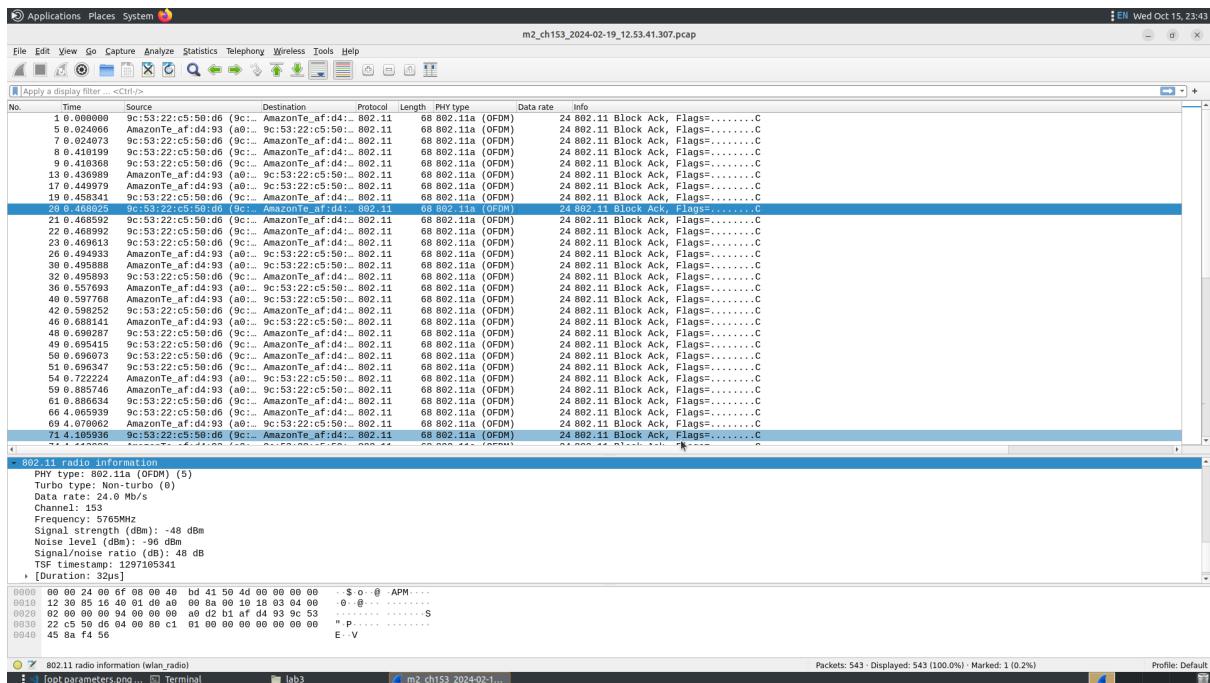


Figure 19: SNR and strength

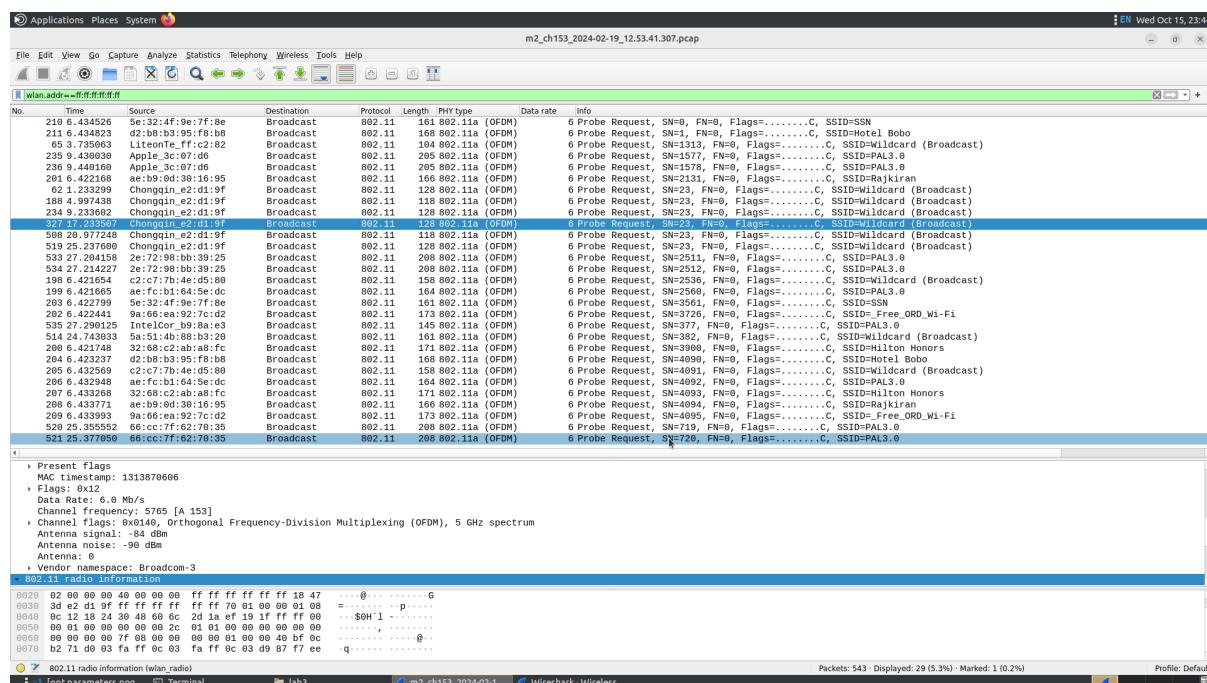


Figure 20: SNR of hightraffic basic service sets and their MAC addresses